

knn

November 25, 2024

```
[7]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Load the dataset
df = pd.read_csv('/content/Zoo.csv') # Replace with the actual file path

# Basic exploration
print("First 5 rows of the dataset:")
print(df.head()) # View first few rows
print("\nDataset Information:")
print(df.info()) # Check for data types and null values
print("\nSummary Statistics:")
print(df.describe()) # Summary statistics for numeric columns

# Pair plot (for numeric features)
try:
    sns.pairplot(df, hue='class') # Replace 'class' with the target column name
    plt.title('Pairwise Feature Relationships')
    plt.show()
except Exception as e:
    print(f"Pair plot error: {e}")

# Class distribution (for categorical target)
try:
    sns.countplot(x='class', data=df) # Replace 'class' with the target column name
    plt.title('Class Distribution')
    plt.show()
except Exception as e:
    print(f"Class distribution plot error: {e}")
```

First 5 rows of the dataset:

	animal name	hair	feathers	eggs	milk	airborne	aquatic	predator	\
0	aardvark	1	0	0	1	0	0	1	
1	antelope	1	0	0	1	0	0	0	
2	bass	0	0	1	0	0	1	1	
3	bear	1	0	0	1	0	0	1	

4	boar	1	0	0	1	0	0	1
---	------	---	---	---	---	---	---	---

	toothed	backbone	breathes	venomous	fins	legs	tail	domestic	catsize	\
0	1	1	1	0	0	4	0	0	1	
1	1	1	1	0	0	4	1	0	1	
2	1	1	0	0	1	0	1	0	0	
3	1	1	1	0	0	4	0	0	1	
4	1	1	1	0	0	4	1	0	1	

	type
0	1
1	1
2	4
3	1
4	1

Dataset Information:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 101 entries, 0 to 100

Data columns (total 18 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	----
0	animal name	101 non-null	object
1	hair	101 non-null	int64
2	feathers	101 non-null	int64
3	eggs	101 non-null	int64
4	milk	101 non-null	int64
5	airborne	101 non-null	int64
6	aquatic	101 non-null	int64
7	predator	101 non-null	int64
8	toothed	101 non-null	int64
9	backbone	101 non-null	int64
10	breathes	101 non-null	int64
11	venomous	101 non-null	int64
12	fins	101 non-null	int64
13	legs	101 non-null	int64
14	tail	101 non-null	int64
15	domestic	101 non-null	int64
16	catsize	101 non-null	int64
17	type	101 non-null	int64

dtypes: int64(17), object(1)

memory usage: 14.3+ KB

None

Summary Statistics:

	hair	feathers	eggs	milk	airborne	aquatic	\
count	101.000000	101.000000	101.000000	101.000000	101.000000	101.000000	
mean	0.425743	0.198020	0.584158	0.405941	0.237624	0.356436	


std	0.496921	0.400495	0.495325	0.493522	0.427750	0.481335
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000
75%	1.000000	0.000000	1.000000	1.000000	0.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

	predator	toothed	backbone	breathes	venomous	fins \
count	101.000000	101.000000	101.000000	101.000000	101.000000	101.000000
mean	0.554455	0.603960	0.821782	0.792079	0.079208	0.168317
std	0.499505	0.491512	0.384605	0.407844	0.271410	0.376013
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	1.000000	1.000000	0.000000	0.000000
50%	1.000000	1.000000	1.000000	1.000000	0.000000	0.000000
75%	1.000000	1.000000	1.000000	1.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

	legs	tail	domestic	catsize	type
count	101.000000	101.000000	101.000000	101.000000	101.000000
mean	2.841584	0.742574	0.128713	0.435644	2.831683
std	2.033385	0.439397	0.336552	0.498314	2.102709
min	0.000000	0.000000	0.000000	0.000000	1.000000
25%	2.000000	0.000000	0.000000	0.000000	1.000000
50%	4.000000	1.000000	0.000000	0.000000	2.000000
75%	4.000000	1.000000	0.000000	1.000000	4.000000
max	8.000000	1.000000	1.000000	1.000000	7.000000

Pair plot error: 'class'

Class distribution plot error: Could not interpret value `class` for `x`. An entry with this name does not appear in `data`.



```
[9]: # Handle missing values
numeric_cols = df.select_dtypes(include=['number']).columns # Get numeric
↳ columns
df[numeric_cols] = df[numeric_cols].fillna(df[numeric_cols].mean()) # Mean
↳ imputation for numeric cols only
# If you have a specific categorical column you want to handle:
# df['categorical_col'].fillna(df['categorical_col'].mode()[0], inplace=True)
↳ # Mode imputation for a specific column

# Handle outliers (using IQR)
Q1 = df[numeric_cols].quantile(0.25)
Q3 = df[numeric_cols].quantile(0.75)
IQR = Q3 - Q1
```

```
df = df[~((df[numeric_cols] < (Q1 - 1.5 * IQR)) | (df[numeric_cols] > (Q3 + 1.5 * IQR))).any(axis=1)] # Remove outliers

print(df.isnull().sum()) # Confirm no missing values remain
```

```
animal name    0
hair           0
feathers       0
eggs           0
milk           0
airborne       0
aquatic        0
predator       0
toothed        0
backbone       0
breathes       0
venomous       0
fins           0
legs           0
tail           0
domestic       0
catsize        0
type           0
dtype: int64
```

```
[11]: # Separate features and target
X = df.drop(columns='animal name') # Assuming 'animal name' is the target
    ↪ column name
y = df['animal name']

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪ random_state=42)

print(f"Training set size: {X_train.shape}")
print(f"Testing set size: {X_test.shape}")
```

```
Training set size: (26, 17)
Testing set size: (7, 17)
```

```
[12]: from sklearn.neighbors import KNeighborsClassifier

# Instantiate the model
k = 5 # Example value for K, can be optimized later
knn = KNeighborsClassifier(n_neighbors=k, metric='euclidean') # Using
    ↪ Euclidean distance
```

```
# Fit the model
knn.fit(X_train, y_train)
```

```
[12]: KNeighborsClassifier(metric='euclidean')
```

```
[13]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report

# Predict on test set
y_pred = knn.predict(X_test)

# Evaluate performance
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted') # Weighted for multi-class
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

```
Accuracy: 0.0
Precision: 0.0
Recall: 0.0
F1 Score: 0.0
```

```
Classification Report:
```

	precision	recall	f1-score	support
antelope	0.00	0.00	0.00	0.0
bear	0.00	0.00	0.00	0.0
boar	0.00	0.00	0.00	0.0
cheetah	0.00	0.00	0.00	0.0
frog	0.00	0.00	0.00	1.0
giraffe	0.00	0.00	0.00	1.0
lion	0.00	0.00	0.00	0.0
mink	0.00	0.00	0.00	1.0
mongoose	0.00	0.00	0.00	1.0
opossum	0.00	0.00	0.00	1.0
squirrel	0.00	0.00	0.00	1.0
wallaby	0.00	0.00	0.00	1.0
accuracy			0.00	7.0
macro avg	0.00	0.00	0.00	7.0

weighted avg 0.00 0.00 0.00 7.0

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in labels
with no true samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in labels
with no true samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in labels
with no true samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Recall is ill-defined and being set to 0.0 in labels
with no true samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
[19]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder

# ... (previous code) ...

# Reduce to 2 features for visualization (if needed)
X_train_vis = X_train.iloc[:, :2]
```

```

X_test_vis = X_test.iloc[:, :2]

# Create a LabelEncoder to convert class labels to numerical values
le = LabelEncoder()
y_train_encoded = le.fit_transform(y_train)

# Train KNN on the reduced dataset using encoded labels
knn.fit(X_train_vis, y_train_encoded)

# Create a meshgrid for decision boundaries
x_min, x_max = X_train_vis.iloc[:, 0].min() - 1, X_train_vis.iloc[:, 0].max() + 1
y_min, y_max = X_train_vis.iloc[:, 1].min() - 1, X_train_vis.iloc[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1), np.arange(y_min, y_max, 0.1))

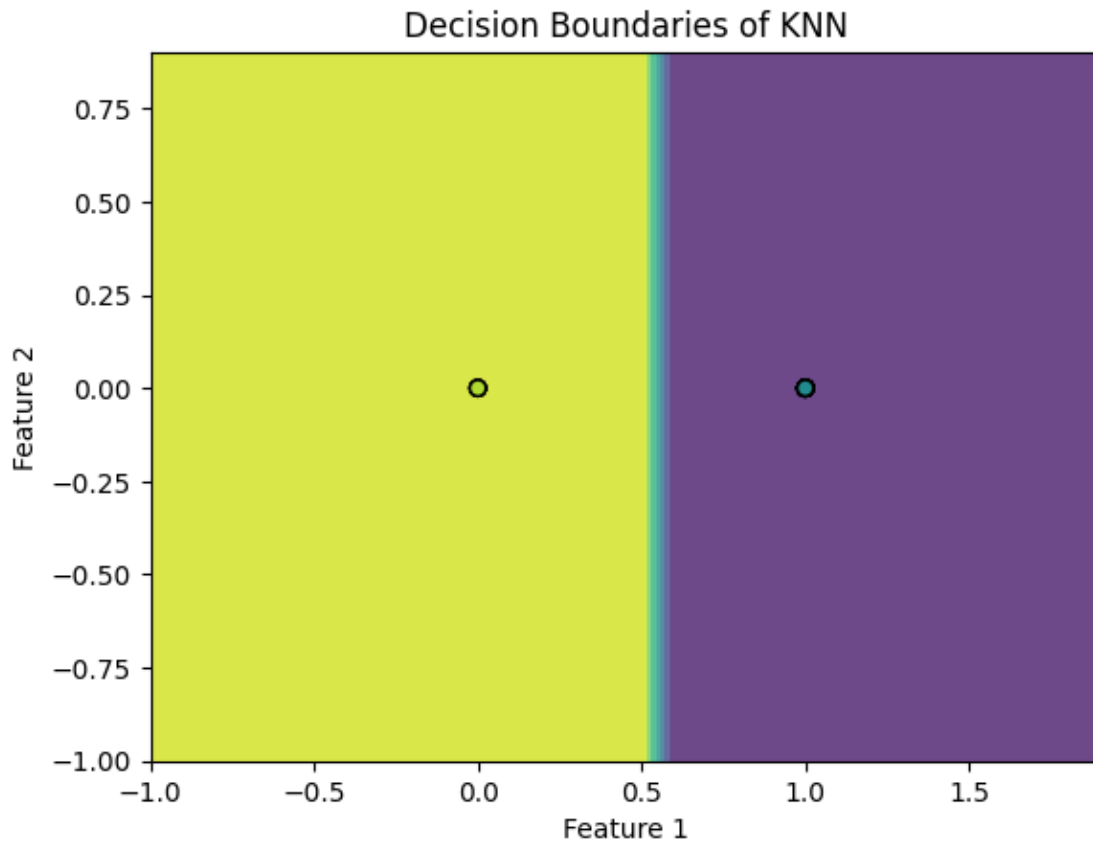
# Predict on meshgrid
Z = knn.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

# Plot decision boundaries
plt.contourf(xx, yy, Z, alpha=0.8)
plt.scatter(X_train_vis.iloc[:, 0], X_train_vis.iloc[:, 1], c=y_train_encoded,
            edgecolor='k', marker='o') # Use encoded labels for plotting
plt.title('Decision Boundaries of KNN')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()

```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:493: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted with feature names

```
warnings.warn(
```

```
[24]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.neighbors import KNeighborsClassifier

# Ensure the dataset has at least 2 features
if X_train.shape[1] < 2 or X_test.shape[1] < 2:
    raise ValueError("The dataset must have at least two features for_
↳ visualization.")

# Select two features for visualization (first two features)
X_train_vis = X_train.iloc[:, :2]
X_test_vis = X_test.iloc[:, :2]

# Train a KNN classifier
try:
    knn_vis = KNeighborsClassifier(n_neighbors=5, metric='euclidean')
    knn_vis.fit(X_train_vis, y_train)
except Exception as e:
    print(f"Error training KNN: {e}")
    raise
```

```

# Create a meshgrid for visualization
x_min, x_max = X_train_vis.iloc[:, 0].min() - 1, X_train_vis.iloc[:, 0].max() + 1
y_min, y_max = X_train_vis.iloc[:, 1].min() - 1, X_train_vis.iloc[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1), np.arange(y_min, y_max, 0.1))

# Predict grid points
try:
    Z = knn_vis.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
except Exception as e:
    print(f"Error predicting meshgrid points: {e}")
    raise

# Plot decision boundaries
plt.figure(figsize=(10, 6))
plt.contourf(xx, yy, Z, alpha=0.8, cmap='coolwarm')

# Plot training points
plt.scatter(X_train_vis.iloc[:, 0], X_train_vis.iloc[:, 1], c=y_train,
            edgecolor='k', marker='o', cmap='coolwarm', label='Training data')

# Plot testing points
plt.scatter(X_test_vis.iloc[:, 0], X_test_vis.iloc[:, 1], c=y_test,
            edgecolor='k', marker='x', cmap='coolwarm', label='Testing data')

plt.title("Decision Boundaries of KNN Classifier")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.legend()
plt.show()

```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:493: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted with feature names

```
warnings.warn(
```

```

-----
TypeError                                Traceback (most recent call last)
<ipython-input-24-e1e6a8b45154> in <cell line: 36>()
    34 # Plot decision boundaries
    35 plt.figure(figsize=(10, 6))
--> 36 plt.contourf(xx, yy, Z, alpha=0.8, cmap='coolwarm')
    37
    38 # Plot training points

```

```

/usr/local/lib/python3.10/dist-packages/matplotlib/pyplot.py in contourf(data,
↳ *args, **kwargs)
    2936 @_copy_docstring_and_deprecators(Axes.contourf)
    2937 def contourf(*args, data=None, **kwargs) -> QuadContourSet:
-> 2938     __ret = gca().contourf(
    2939         *args, **({"data": data} if data is not None else {}), **kwargs
    2940     )

/usr/local/lib/python3.10/dist-packages/matplotlib/__init__.py in inner(ax,
↳ data, *args, **kwargs)
    1463     def inner(ax, *args, data=None, **kwargs):
    1464         if data is None:
-> 1465             return func(ax, *map(sanitize_sequence, args), **kwargs)
    1466
    1467         bound = new_sig.bind(ax, *args, **kwargs)

/usr/local/lib/python3.10/dist-packages/matplotlib/axes/_axes.py in
↳ contourf(self, *args, **kwargs)
    6526         """
    6527         kwargs['filled'] = True
-> 6528         contours = mcontour.QuadContourSet(self, *args, **kwargs)
    6529         self._request_autoscale_view()
    6530         return contours

/usr/local/lib/python3.10/dist-packages/matplotlib/contour.py in __init__(self,
↳ ax, levels, filled, linewidths, linestyles, hatches, alpha, origin, extent,
↳ cmap, colors, norm, vmin, vmax, extend, antialiased, nchunk, locator,
↳ transform, negative_linestyles, clip_path, *args, **kwargs)
    845         mpl.rcParams['contour.negative_linestyle']
    846
--> 847         kwargs = self._process_args(*args, **kwargs)
    848         self._process_levels()
    849

/usr/local/lib/python3.10/dist-packages/matplotlib/contour.py in
↳ _process_args(self, corner_mask, algorithm, *args, **kwargs)
    1531         self._corner_mask = corner_mask
    1532
-> 1533         x, y, z = self._contour_args(args, kwargs)
    1534
    1535         contour_generator = contourpy.contour_generator(

/usr/local/lib/python3.10/dist-packages/matplotlib/contour.py in
↳ _contour_args(self, args, kwargs)
    1575         else:
    1576             raise _api.nargs_error(fn, takes="from 1 to 4", given=nargs
-> 1577         z = ma.masked_invalid(z, copy=False)

```


```

1578         self.zmax = z.max().astype(float)
1579         self.zmin = z.min().astype(float)

/usr/local/lib/python3.10/dist-packages/numpy/ma/core.py in masked_invalid(a,
↳copy)
2358         """
2359         a = np.array(a, copy=False, subok=True)
-> 2360         res = masked_where(~(np.isfinite(a)), a, copy=copy)
2361         # masked_invalid previously never returned nomask as a mask and
↳doing so
2362         # threw off matplotlib (gh-22842). So use shrink=False:

TypeError: ufunc 'isfinite' not supported for the input types, and the inputs
↳could not be safely coerced to any supported types according to the casting
↳rule 'safe'

```



```

[25]: print(X_train.dtypes) # Check data types of features
      print(y_train.dtype)  # Check data type of target

```

```

hair          int64
feathers       int64
eggs          int64
milk          int64
airborne      int64
aquatic       int64

```

```

predator    int64
toothed     int64
backbone    int64
breathes    int64
venomous    int64
fins        int64
legs        int64
tail        int64
domestic    int64
catsize     int64
type        int64
dtype: object
object

```

```

[29]: X_train = X_train.apply(pd.to_numeric, errors='coerce')
      X_test = X_test.apply(pd.to_numeric, errors='coerce')

```

```

[30]: import numpy as np
      import matplotlib.pyplot as plt
      from sklearn.neighbors import KNeighborsClassifier

      # Select the first two features
      X_train_vis = X_train.iloc[:, :2]
      X_test_vis = X_test.iloc[:, :2]

      # Train KNN classifier
      knn_vis = KNeighborsClassifier(n_neighbors=5, metric='euclidean')
      knn_vis.fit(X_train_vis, y_train)

      # Create meshgrid
      x_min, x_max = X_train_vis.iloc[:, 0].min() - 1, X_train_vis.iloc[:, 0].max() + 1
      y_min, y_max = X_train_vis.iloc[:, 1].min() - 1, X_train_vis.iloc[:, 1].max() + 1
      xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1), np.arange(y_min, y_max, 0.1))

      # Predict on meshgrid
      Z = knn_vis.predict(np.c_[xx.ravel(), yy.ravel()])
      Z = Z.reshape(xx.shape)

      # Plot decision boundaries
      plt.figure(figsize=(10, 6))
      plt.contourf(xx, yy, Z, alpha=0.8, cmap='coolwarm')

      # Plot training points
      plt.scatter(X_train_vis.iloc[:, 0], X_train_vis.iloc[:, 1], c=y_train,
                  edgecolor='k', marker='o', cmap='coolwarm', label='Training data')

```

```

# Plot testing points
plt.scatter(X_test_vis.iloc[:, 0], X_test_vis.iloc[:, 1], c=y_test,
            edgecolor='k', marker='x', cmap='coolwarm', label='Testing data')

plt.title("Decision Boundaries of KNN Classifier")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.legend()
plt.show()

```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:493: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted with feature names

```
warnings.warn(
```

```

-----
ValueError                                Traceback (most recent call last)
/usr/local/lib/python3.10/dist-packages/matplotlib/axes/_axes.py in
    ↪ _parse_scatter_color_args(c, edgecolors, kwargs, xsize, get_next_color_func)
    4482                 try: # Is 'c' acceptable as PathCollection facecolors?
--> 4483                     colors = mcolors.to_rgba_array(c)
    4484             except (TypeError, ValueError) as err:

/usr/local/lib/python3.10/dist-packages/matplotlib/colors.py in to_rgba_array(c,
    ↪ alpha)
    504     else:
--> 505         rgba = np.array([to_rgba(cc) for cc in c])
    506

/usr/local/lib/python3.10/dist-packages/matplotlib/colors.py in <listcomp>(.0)
    504     else:
--> 505         rgba = np.array([to_rgba(cc) for cc in c])
    506

/usr/local/lib/python3.10/dist-packages/matplotlib/colors.py in to_rgba(c, alph.)
    301     if rgba is None: # Suppress exception chaining of cache lookup
    ↪ failure.
--> 302         rgba = _to_rgba_no_colorcycle(c, alpha)
    303         try:

/usr/local/lib/python3.10/dist-packages/matplotlib/colors.py in
    ↪ _to_rgba_no_colorcycle(c, alpha)
    383             return c, c, c, alpha if alpha is not None else 1.
--> 384             raise ValueError(f"Invalid RGBA argument: {orig_c!r}")
    385         # turn 2-D array into 1-D array

```

```
ValueError: Invalid RGBA argument: 'wallaby'
```

The above exception was the direct cause of the following exception:

```
ValueError                                Traceback (most recent call last)
<ipython-input-30-a494d8a7aa1b> in <cell line: 30>()
    28
    29 # Plot testing points
--> 30 plt.scatter(X_test_vis.iloc[:, 0], X_test_vis.iloc[:, 1], c=y_test,
    ↪ edgecolor='k', marker='x', cmap='coolwarm', label='Testing data')
    31
    32 plt.title("Decision Boundaries of KNN Classifier")

/usr/local/lib/python3.10/dist-packages/matplotlib/pyplot.py in scatter(x, y, s,
    ↪ c, marker, cmap, norm, vmin, vmax, alpha, linewidths, edgecolors,
    ↪ plotnonfinite, data, **kwargs)
    3685     **kwargs,
    3686 ) -> PathCollection:
-> 3687     __ret = gca().scatter(
    3688         x,
    3689         y,

/usr/local/lib/python3.10/dist-packages/matplotlib/_init__.py in inner(ax,
    ↪ data, *args, **kwargs)
    1463     def inner(ax, *args, data=None, **kwargs):
    1464         if data is None:
-> 1465             return func(ax, *map(sanitize_sequence, args), **kwargs)
    1466
    1467         bound = new_sig.bind(ax, *args, **kwargs)

/usr/local/lib/python3.10/dist-packages/matplotlib/axes/_axes.py in
    ↪ scatter(self, x, y, s, c, marker, cmap, norm, vmin, vmax, alpha, linewidths,
    ↪ edgecolors, plotnonfinite, **kwargs)
    4668         orig_edgecolor = kwargs.get('edgecolor', None)
    4669         c, colors, edgecolors = \
-> 4670         self._parse_scatter_color_args(
    4671             c, edgecolors, kwargs, x.size,
    4672             get_next_color_func=self._get_patches_for_fill.
    ↪ get_next_color)


/usr/local/lib/python3.10/dist-packages/matplotlib/axes/_axes.py in
    ↪ _parse_scatter_color_args(c, edgecolors, kwargs, xsize, get_next_color_func)
    4490         # Both the mapping *and* the RGBA conversion failed
    ↪ pretty
    4491         # severe failure => one may appreciate a verbose
    ↪ feedback.
```

```

-> 4492                                     raise ValueError(
    4493                                     f"'c' argument must be a color, a sequence of
↪ colors, "
    4494                                     f"or a sequence of numbers, not {c!r}") from er:

ValueError: 'c' argument must be a color, a sequence of colors, or a sequence of
↪ numbers, not 96      wallaby
48      mink
84      squirrel
50      mongoose
25      frog
28      giraffe
54      opossum
Name: animal name, dtype: object

```



```

[31]: from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()
y_train = le.fit_transform(y_train)  # Encode training labels
y_test = le.transform(y_test)        # Encode testing labels

```

```

-----
ValueError                                Traceback (most recent call last)
<ipython-input-31-8af8943cd32f> in <cell line: 5>()
    3 le = LabelEncoder()

```



```

4 y_train = le.fit_transform(y_train) # Encode training labels
----> 5 y_test = le.transform(y_test) # Encode testing labels

/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_label.py in
-> transform(self, y)
    130         """
    131         check_is_fitted(self)
--> 132         y = column_or_1d(y, dtype=self.classes_.dtype, warn=True)
    133         # transform of empty array is empty array
    134         if _num_samples(y) == 0:

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py in
-> column_or_1d(y, dtype, warn)
    1379         """
    1380         xp, _ = get_namespace(y)
-> 1381         y = check_array(
    1382             y,
    1383             ensure_2d=False,

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py in
-> check_array(array, accept_sparse, accept_large_sparse, dtype, order, copy,
-> force_writeable, force_all_finite, ensure_2d, allow_nd, ensure_min_samples,
-> ensure_min_features, estimator, input_name)
    1008             input_name=input_name,
    1009         )
-> 1010         array = xp.astype(array, dtype, copy=False)
    1011         else:
    1012             array = _asarray_with_order(array, order=order,
-> dtype=dtype, xp=xp)

/usr/local/lib/python3.10/dist-packages/sklearn/utils/_array_api.py in
-> astype(self, x, dtype, copy, casting)
    388     def astype(self, x, dtype, *, copy=True, casting="unsafe"):
    389         # astype is not defined in the top level NumPy namespace
--> 390         return x.astype(dtype, copy=copy, casting=casting)
    391
    392     def asarray(self, x, *, dtype=None, device=None, copy=None): # noqa

ValueError: invalid literal for int() with base 10: 'wallaby'

```

```

[32]: print("y_train type:", type(y_train))
print("Unique values in y_train:", y_train.unique() if hasattr(y_train,
-> "unique") else np.unique(y_train))

```

```

y_train type: <class 'numpy.ndarray'>
Unique values in y_train: [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17
18 19 20 21 22 23]

```

24 25]

```
[33]: y_train = y_train.astype(str)
      y_test = y_test.astype(str)
```

```
[37]: import pandas as pd
      from sklearn.preprocessing import LabelEncoder

      # Ensure y_train and y_test are strings, which is common for categorical labels
      y_train = y_train.astype(str)
      y_test = y_test.astype(str)

      # Combine y_train and y_test into one Series
      combined_labels = pd.concat([y_train, y_test], axis=0)

      # Initialize and fit the LabelEncoder on combined labels
      le = LabelEncoder()
      le.fit(combined_labels)

      # Transform y_train and y_test using the learned encoder
      y_train = le.transform(y_train)
      y_test = le.transform(y_test)

      # Print mapping for reference
      print("Label mapping:", dict(zip(le.classes_, le.transform(le.classes_))))
```

```
Label mapping: {'0': 0, '1': 1, '10': 2, '11': 3, '12': 4, '13': 5, '14': 6,
'15': 7, '16': 8, '17': 9, '18': 10, '19': 11, '2': 12, '20': 13, '21': 14,
'22': 15, '23': 16, '24': 17, '25': 18, '3': 19, '4': 20, '5': 21, '6': 22, '7':
23, '8': 24, '9': 25, 'frog': 26, 'giraffe': 27, 'mink': 28, 'mongoose': 29,
'opossum': 30, 'squirrel': 31, 'wallaby': 32}
```

```
[38]: print("Data types after conversion:", y_train.dtype, y_test.dtype)
      print("Encoded y_train:", y_train)
      print("Encoded y_test:", y_test)
```

```
Data types after conversion: int64 int64
Encoded y_train: [24 12  0 13 21 14 19 18  3  1  4  8 11  6  2  9 25 22  5 15 23
10 16 17
20  7]
Encoded y_test: [32 28 31 29 26 27 30]
```

```
[40]: import numpy as np
      import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.preprocessing import LabelEncoder
      import matplotlib.pyplot as plt
```

```

# Assuming you have your dataset loaded into X and y
# Example: Replace this with your actual dataset
# X = your_features_dataframe
# y = your_target_series

# Ensure y_train and y_test are pandas Series
y_train = pd.Series(y_train)
y_test = pd.Series(y_test)

# Combine y_train and y_test into one Series for fitting LabelEncoder
combined_labels = pd.concat([y_train, y_test], axis=0)

# Initialize LabelEncoder and fit it on the combined labels
le = LabelEncoder()
le.fit(combined_labels)

# Transform y_train and y_test using the learned encoder
y_train = le.transform(y_train)
y_test = le.transform(y_test)

# Print the label mapping for reference
print("Label mapping:", dict(zip(le.classes_, le.transform(le.classes_))))

# Select the first two features for visualization (ensure your dataset has at
↳ least 2 features)
X_train_vis = X_train.iloc[:, :2]
X_test_vis = X_test.iloc[:, :2]

# Train the KNN classifier (with k=5 and Euclidean distance)
knn_vis = KNeighborsClassifier(n_neighbors=5, metric='euclidean')
knn_vis.fit(X_train_vis, y_train)

# Create a meshgrid for plotting the decision boundaries
x_min, x_max = X_train_vis.iloc[:, 0].min() - 1, X_train_vis.iloc[:, 0].max() +
↳ 1
y_min, y_max = X_train_vis.iloc[:, 1].min() - 1, X_train_vis.iloc[:, 1].max() +
↳ 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1), np.arange(y_min, y_max, 0.1))

# Predict on the meshgrid to plot decision boundaries
Z = knn_vis.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

# Plot the decision boundaries
plt.figure(figsize=(10, 6))
plt.contourf(xx, yy, Z, alpha=0.8, cmap='coolwarm')

```

```

# Plot the training points
plt.scatter(X_train_vis.iloc[:, 0], X_train_vis.iloc[:, 1], c=y_train,
            edgecolor='k', marker='o', cmap='coolwarm', label='Training data')

# Plot the testing points
plt.scatter(X_test_vis.iloc[:, 0], X_test_vis.iloc[:, 1], c=y_test,
            edgecolor='k', marker='x', cmap='coolwarm', label='Testing data')

plt.title("Decision Boundaries of KNN Classifier")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.legend()
plt.show()

```

Label mapping: {'0': 0, '1': 1, '10': 2, '11': 3, '12': 4, '13': 5, '14': 6, '15': 7, '16': 8, '17': 9, '18': 10, '19': 11, '2': 12, '20': 13, '21': 14, '22': 15, '23': 16, '24': 17, '25': 18, '26': 19, '27': 20, '28': 21, '29': 22, '3': 23, '30': 24, '31': 25, '32': 26, '4': 27, '5': 28, '6': 29, '7': 30, '8': 31, '9': 32}

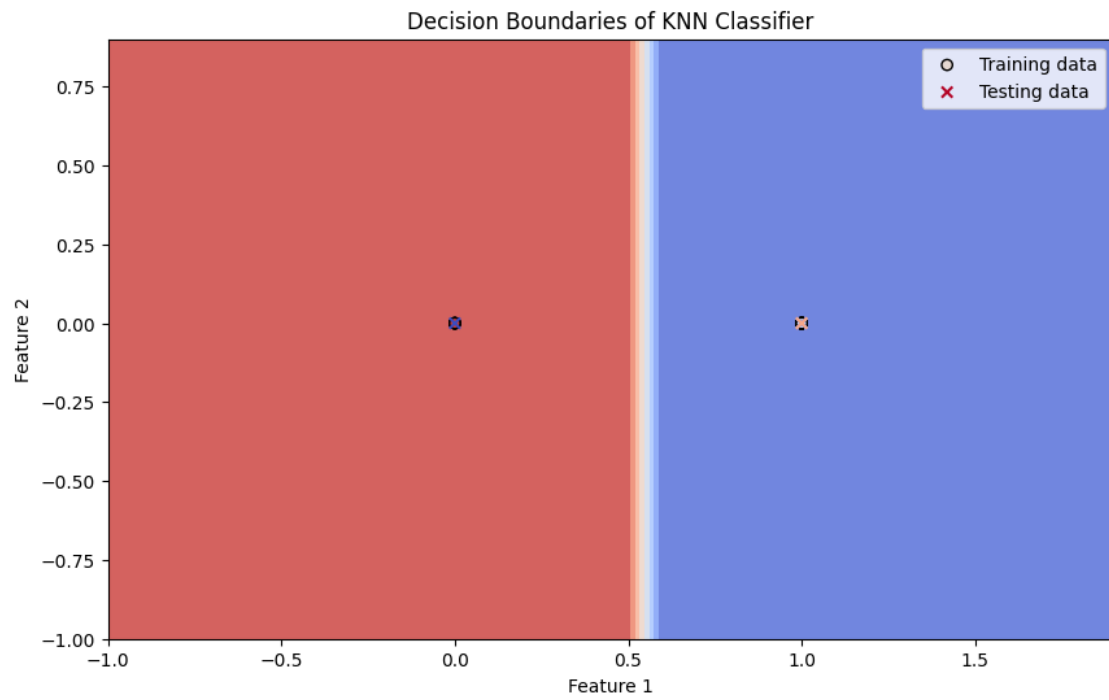
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:493: UserWarning: X does not have valid feature names, but KNeighborsClassifier was fitted with feature names

```
warnings.warn(
<ipython-input-40-ffffef41ec01f>:56: UserWarning: You passed a
edgecolor/edgecolors ('k') for an unfilled marker ('x'). Matplotlib is ignoring
the edgecolor in favor of the facecolor. This behavior may change in the
future.
```

```

plt.scatter(X_test_vis.iloc[:, 0], X_test_vis.iloc[:, 1], c=y_test,
            edgecolor='k', marker='x', cmap='coolwarm', label='Testing data')

```



[]: