

## SVM

November 25, 2024

```
[1]: import pandas as pd
df = pd.read_csv('/content/mushroom.csv')
```

```
[3]: print(df.head())
print(df.info())
print(df.describe())
```

	Unnamed: 0	cap_shape	cap_surface	cap_color	bruises	odor	gill_attachment	\
0	1167	sunken	scaly	white	no	anise	descending	
1	1037	sunken	fibrous	red	no	anise	notched	
2	309	flat	grooves	purple	yes	foul	descending	
3	282	bell	scaly	pink	yes	fishy	notched	
4	820	flat	smooth	yellow	yes	musty	free	

	gill_spacing	gill_size	gill_color	...	veil_type	veil_color	ring_number	\
0	distant	broad	pink	...	partial	brown	two	
1	crowded	narrow	chocolate	...	universal	brown	two	
2	crowded	broad	purple	...	universal	yellow	two	
3	close	broad	orange	...	partial	yellow	two	
4	crowded	narrow	orange	...	universal	white	none	

	ring_type	spore_print_color	population	habitat	class	stalk_height	\
0	sheathing	chocolate	clustered	waste	poisonous	14.276173	
1	sheathing	brown	numerous	waste	edible	3.952715	
2	sheathing	purple	abundant	waste	poisonous	9.054265	
3	cobwebby	green	clustered	grasses	poisonous	5.226499	
4	none	yellow	clustered	urban	poisonous	14.037532	

	cap_diameter
0	5.054983
1	19.068319
2	7.205884
3	20.932692
4	12.545245

[5 rows x 26 columns]

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
```

Data columns (total 26 columns):

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	2000 non-null	int64
1	cap_shape	2000 non-null	object
2	cap_surface	2000 non-null	object
3	cap_color	2000 non-null	object
4	bruises	2000 non-null	object
5	odor	2000 non-null	object
6	gill_attachment	2000 non-null	object
7	gill_spacing	2000 non-null	object
8	gill_size	2000 non-null	object
9	gill_color	2000 non-null	object
10	stalk_shape	2000 non-null	object
11	stalk_root	2000 non-null	object
12	stalk_surface_above_ring	2000 non-null	object
13	stalk_surface_below_ring	2000 non-null	object
14	stalk_color_above_ring	2000 non-null	object
15	stalk_color_below_ring	2000 non-null	object
16	veil_type	2000 non-null	object
17	veil_color	2000 non-null	object
18	ring_number	2000 non-null	object
19	ring_type	2000 non-null	object
20	spore_print_color	2000 non-null	object
21	population	2000 non-null	object
22	habitat	2000 non-null	object
23	class	2000 non-null	object
24	stalk_height	2000 non-null	float64
25	cap_diameter	2000 non-null	float64

dtypes: float64(2), int64(1), object(23)

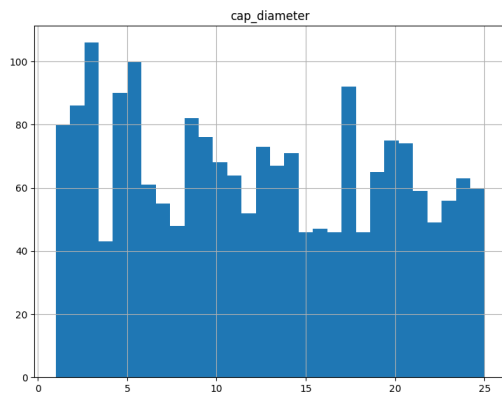
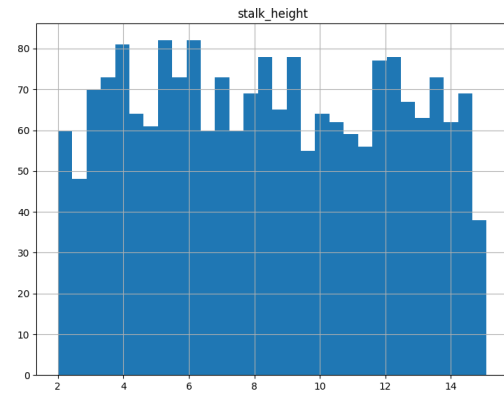

memory usage: 406.4+ KB

None

	Unnamed: 0	stalk_height	cap_diameter
count	2000.000000	2000.000000	2000.000000
mean	624.974000	8.449118	12.314345
std	375.091938	3.697217	7.048845
min	0.000000	2.000000	1.000000
25%	290.000000	5.291009	5.723521
50%	607.000000	8.318596	12.124902
75%	957.250000	11.781272	18.698605
max	1279.000000	15.095066	25.000054

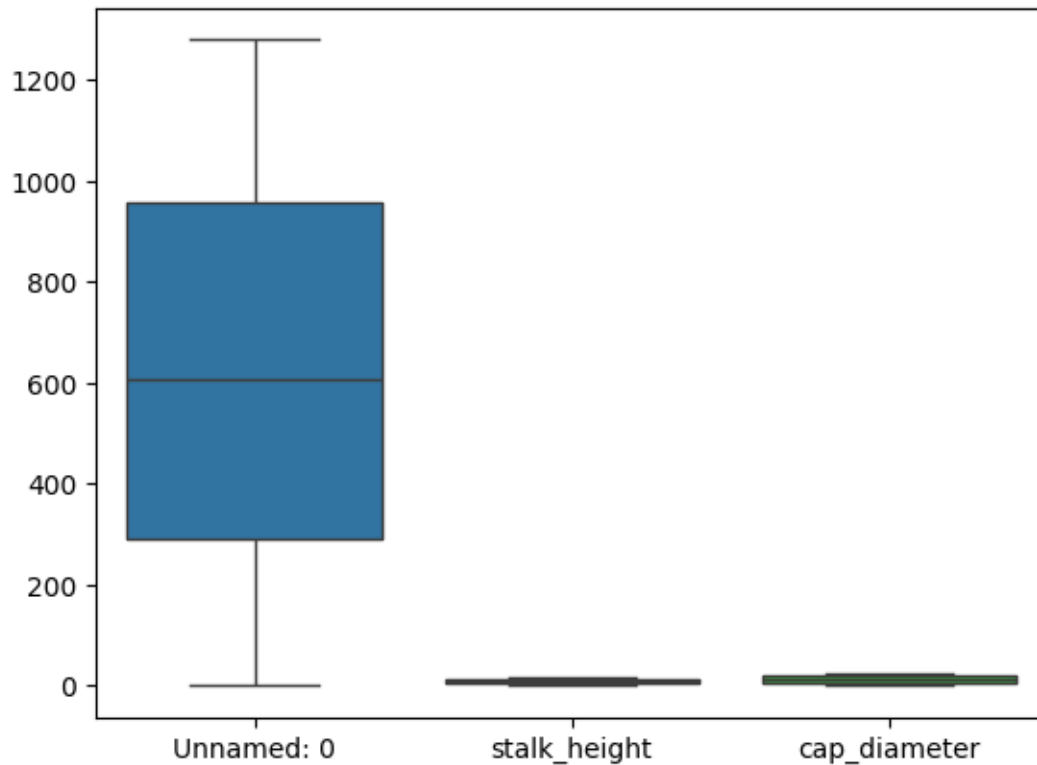
```
[4]: df.hist(bins=30, figsize=(20, 15))
```

```
[4]: array([[<Axes: title={'center': 'Unnamed: 0'}>,  
          <Axes: title={'center': 'stalk_height'}>],  
          [<Axes: title={'center': 'cap_diameter'}>, <Axes: >]], dtype=object)
```



```
[5]: import seaborn as sns
sns.boxplot(data=df)
```

```
[5]: <Axes: >
```




```
[6]: df.plot(kind='density', subplots=True, layout=(8, 8), sharex=False,
      figsize=(20, 20))
```

[illegible]

```

    <Axes: ylabel='Density'>, <Axes: ylabel='Density'>],
    [<Axes: ylabel='Density'>, <Axes: ylabel='Density'>,
    <Axes: ylabel='Density'>, <Axes: ylabel='Density'>,
    <Axes: ylabel='Density'>, <Axes: ylabel='Density'>,
    <Axes: ylabel='Density'>, <Axes: ylabel='Density'>],
    [<Axes: ylabel='Density'>, <Axes: ylabel='Density'>,
    <Axes: ylabel='Density'>, <Axes: ylabel='Density'>,
    <Axes: ylabel='Density'>, <Axes: ylabel='Density'>,
    <Axes: ylabel='Density'>, <Axes: ylabel='Density'>],
    [<Axes: ylabel='Density'>, <Axes: ylabel='Density'>,
    <Axes: ylabel='Density'>, <Axes: ylabel='Density'>,
    <Axes: ylabel='Density'>, <Axes: ylabel='Density'>,
    <Axes: ylabel='Density'>, <Axes: ylabel='Density'>]], dtype=object)

```



```

[11]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

```

```

[12]: print(df.isnull().sum())

```

```

Unnamed: 0          0
cap_shape          2000
cap_surface        2000
cap_color           0
bruises            0
odor               0
gill_attachment    0
gill_spacing       0
gill_size          0
gill_color         0
stalk_shape        0
stalk_root         0
stalk_surface_above_ring  0
stalk_surface_below_ring  0

```

```
stalk_color_above_ring      0
stalk_color_below_ring      0
veil_type                   0
veil_color                   0
ring_number                  0
ring_type                    0
spore_print_color            0
population                   0
habitat                      0
class                        0
stalk_height                 0
cap_diameter                 0
dtype: int64
```

```
[14]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np # Import numpy with the alias np

df_numeric = df.select_dtypes(include=[np.number])
corr_matrix = df_numeric.corr()
```


```
[15]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

# Load your dataset
df = pd.read_csv('/content/mushroom.csv')

# Select only numeric columns if necessary
df_numeric = df.select_dtypes(include=[np.number])

# Calculate the correlation matrix
corr_matrix = df_numeric.corr()

# Plot the heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True)
plt.show()
```



```
[16]: from sklearn.preprocessing import LabelEncoder
df = df.apply(LabelEncoder().fit_transform)
```

```
[18]: from sklearn.model_selection import train_test_split


# Replace 'class' with the actual target column name if it's different
X = df.drop('class', axis=1)
y = df['class']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

```
[20]: sns.pairplot(df, hue='class')
```

Output hidden; open in <https://colab.research.google.com> to view.

```
[21]: sns.countplot(x='class', data=df)
```

```
[21]: <Axes: xlabel='class', ylabel='count'>
```



```
[22]: from sklearn.svm import SVC
      model = SVC()
      model.fit(X_train, y_train)
```

[22]: SVC()

```
[23]: model.fit(X_train, y_train)
```

[23]: SVC()

```
[24]: from sklearn.metrics import classification_report
      y_pred = model.predict(X_test)
      print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	133
1	0.67	1.00	0.80	267
accuracy			0.67	400
macro avg	0.33	0.50	0.40	400



weighted avg	0.45	0.67	0.53	400
--------------	------	------	------	-----

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
[25]: # For 2D visualization, we can plot decision boundaries, for example
```

```
[26]: model = SVC(kernel='linear', C=1.0)
      model.fit(X_train, y_train)
```

```
[26]: SVC(kernel='linear')
```

```
[27]: kernels = ['linear', 'poly', 'rbf']
      for kernel in kernels:
          model = SVC(kernel=kernel)
          model.fit(X_train, y_train)
          y_pred = model.predict(X_test)
          print(f"Kernel: {kernel}")
          print(classification_report(y_test, y_pred))
```

```
Kernel: linear
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	133
1	0.67	1.00	0.80	267
accuracy			0.67	400
macro avg	0.33	0.50	0.40	400
weighted avg	0.45	0.67	0.53	400

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.
```

```

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.

```

```

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.

```

```

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Kernel: poly

	precision	recall	f1-score	support
0	0.00	0.00	0.00	133
1	0.67	1.00	0.80	267
accuracy			0.67	400
macro avg	0.33	0.50	0.40	400
weighted avg	0.45	0.67	0.53	400

```

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.

```

```

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.

```

```

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.

```

```

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Kernel: rbf

	precision	recall	f1-score	support
0	0.00	0.00	0.00	133
1	0.67	1.00	0.80	267
accuracy			0.67	400
macro avg	0.33	0.50	0.40	400
weighted avg	0.45	0.67	0.53	400

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.
```

```
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.
```

```
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531:
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control this
behavior.
```

```
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
[ ]:
```